ABBA6 Implicit

One symmetric projection around a Yoshida composition

Historical model. ABBA6Implicit, its exports and its dedicated study were removed on 2026-10-25. The remaining text describes the archived implementation, not a supported method. Saved results retain their original method and projection labels.

Explicit palindromic operations

All ABBA and BM4 presets inherit `PalindromicMethod`. For signed $d = ch$, the elementary operations are

$$A_d(u, v, t) = (u + df(t, v), v, t), \qquad B_d(u, v, t) = (u, v + df(t, u), t), \qquad T_d(u, v, t) = (u, v, t + d).$$

Only T advances time. A coefficient $c = 1/2$ means duration $h/2$; there is no additional half inside A or B. The optional exact C flow rotates the copy difference and requires an independent coupling object. Both BM4 presets omit C.

In chronological order, standard ABBA2 is $A_{h/2}, B_{h/2}, T_h, B_{h/2}, A_{h/2}$. The distinct symmetric variant $T_{h/2}, A_{h/2}, B_{h/2}, B_{h/2}, A_{h/2}, T_{h/2}$ evaluates every spatial shear at the midpoint time. Both exhibit second order on the tested nonautonomous rotation. Arbitrary palindromes do not inherit fourth or sixth order, and an arithmetic projection need not be reversible or symplectic.

`Composition` denotes an explicit tuple of operations. Consecutive time translations combine as $T_a T_b = T_{a+b}$. The generic `PalindromicImplicit` and `PalindromicMidpoint` methods select one exterior Hairer projection or a final arithmetic mean. The signed spatial flows and passive momentum normalization are retained; observers reconstruct their groups from elementary traces without replay. NumPy and JAX execute the same explicit operations and differentiate their retained shear inputs. See `docs/models/palindromic/simulation/palindromic-operations.md`.

1 Unprojected seven-pair composition

Let $S_{q,t}$ denote the symmetric second-order ABBA map on two spatial copies, with signed duration q and start time t. Set

$$a = 0.78451361047755726382, \qquad b = 0.23557321335935813368,$$
$$c = -1.17767998417887100695, \qquad d = 1.31518632068391121889,$$

and $(w_1, \ldots, w_7) = (a, b, c, d, c, b, a)$. For a trial multiplier μ, the complete base map is evaluated through

$$Y^{(0)} = Ez_n + N\mu, \qquad t_j = t_n + h\sum_{\ell<j} w_\ell, \qquad Y^{(j)} = S_{w_j h, t_j}(Y^{(j-1)}), \qquad \Psi_{h,t_n}(Y^{(0)}) = Y^{(7)}.$$

Both copies remain independent between pairs. Each pair applies an adjoint and a direct map with weights $w_j/2$, giving fourteen stages in total. The weights satisfy, up to round-off,

$$\sum_j w_j = 1, \qquad \sum_j w_j^3 = 0, \qquad \sum_j w_j^5 = 0.$$

These moments and palindromy are useful checks, but are not sufficient on their own to prove order six: the Yoshida solution also satisfies the mixed commutator order condition for a symmetric second-order base.

2 One outer symmetric projection

For $m = 2N_p$ physical coordinates, with N_p independent particles, define $Ez = (z, z)$, $G = [I, -I]$, $N = G^{\mathsf{T}}$ and $P = \frac{1}{2}[I, I]$. One complete step solves

$$r(\mu) = G\Psi_{h,t_n}(Ez_n + N\mu) + 2\mu = 0,$$
$$z_{n+1} = P\big(\Psi_{h,t_n}(Ez_n + N\mu) + N\mu\big).$$

The reduced formulation has m unknowns. The simultaneous formulation solves for (Y_+, μ) using

$$Y_+ - N\mu - \Psi_{h,t_n}(Ez_n + N\mu) = 0, \qquad GY_+ = 0.$$

Both formulations define the same local map and permit Newton or Broyden. There is one multiplier and one nonlinear solve per complete step. This is exactly the ABBA4 architecture with a different unprojected recipe. It intentionally replaces the historical composition of seven projected ABBA2 maps; previously saved trajectories and costs belong to that former map.

3 Signed times, order and geometry

The third and fifth pair durations are negative. Internal times can leave $[t_n, t_n + h]$, reaching approximately $t_n - 0.15759h$ and $t_n + 1.15759h$; the dynamics must be defined throughout the signed path. Adjoint stages evaluate at their starting clock, direct stages at their ending clock.

For a smooth field, the Yoshida base has order six on doubled space. Its exact flow preserves the diagonal, so the local diagonal defect is $O(h^7)$. At $h = 0$, $D_\mu r = 4I$; the implicit function theorem therefore gives the local multiplier $\mu = O(h^7)$. The outer projection preserves order six. The symmetric root construction is reversible and, for the canonical Hamiltonian splitting, symplectic on the physical diagonal. These statements concern the local exact root. Finite stopping tolerances, round-off and nonsmooth fields can limit the observed order and geometry.

4 Accepted traces and passive energy

Only the accepted root contributes to passive energy. The 28 signed shear sources give

$$\Delta\kappa = \frac{1}{2}\sum_j s_j[-\partial_t H(t_j, z_j)].$$

Time and momentum do not enter the nonlinear equation. The ABBA6 observation contains one multiplier, complete-step work metrics and seven unprojected pair snapshots. Work histories have shape $(n_{\text{steps}}, 1)$ rather than $(n_{\text{steps}}, 7)$. Differentiation composes the unprojected tangents first, then differentiates the single outer root.

5 Implementation correspondence and verification

The public class is in `src/methods/extended/abba.py`; it inherits the same initialization, advance and placement validation as ABBA4. Recipes and weights live in `src/methods/extended/abba.py`. The path is `advance`, `solve_projection`, `compose`, then explicit $A/B/T/C$ operations. Common Newton/Broyden services, particle-block tangents and energy quadrature are shared with BM4. Public method names remain unchanged; no forwarding modules preserve retired internal paths. Event conversion lives in `extended/observations.py`.

Tests check the one-solve contract, fourteen signed stages, pair continuity, nonautonomous sixth-order convergence, both residual formulations with both solvers, reversibility, analytic tangents and

symplecticity. The nonlinear convergence test uses an independently tighter DOP853 reference. Current architecture and diagrams are documented in `docs/models/palindromic/simulation/palindromic-simulation-architecture.md`.

6 Physical accepted states and passive energy tracking

For N particles with d physical components, every accepted snapshot contains $Z \in \mathbb{R}^{dN}$ and one scalar time t. An optional energy observer owns separate physical momentum $\kappa \in \mathbb{R}^N$. Thus GC accepted vectors have $2N$ components and FC vectors have $4N$, whether energy is tracked or not. A doubled method prepares temporary $(u, v) \in \mathbb{R}^{2dN}$ copies, allows them to separate, and projects to one physical result before acceptance. The representation stores only the physical size and never owns a clock, momentum, physical coordinate declaration, dynamics or solver.

The normalized passive balance is

$$\dot{\kappa} = -\partial_t H(t, Z), \qquad \kappa(t_0) = 0, \qquad \Delta K(t) = H(t, Z(t)) + \kappa(t) - H(t_0, Z_0).$$

GC split-map contributions are summed in temporary work and divided by two to obtain physical momentum. Tracking leaves physical stages, nonlinear stopping scales and adaptive acceptance unchanged. An autonomous Hamiltonian supplies an explicit zero momentum derivative.

Hairer's constraint is solely $u - v = 0$: the reduced multiplier has $2N$ entries and ABBA's simultaneous spatial solve has $6N$ unknowns. Neither t nor κ is projected. The former eight-component projection is retired. Dynamics declare component order and position/velocity roles through immutable structural data; initial conditions and results retain those data independently. Execution injects $f(t, Z)$ and optional physical derivatives into a fresh method run. Methods own physical stages and projection equations, without retaining the dynamics object. Nonlinear tolerances, iteration limits, differentiation and iteration callbacks belong to the configured algebraic solver. Generic doubled copy operations live in `src/formulations/state.py`; classical, HBVM and adaptive methods need no state formulation.

Methods optionally return retained stage quadrature or a dense interpolant. The energy observer integrates momentum independently, using Hamiltonian functions supplied by execution. Controllers keep only physical accepted states; compiled JAX loops carry the observer accumulator separately. Observers receive physical states. Energy histories in `States.diagnostics` retain their existing labels; `extended_time` has shape (N, n_{samples}) and is broadcast from scalar sample times. Fixed shadow samples leave accepted states and counters unchanged. Adaptive quadrature uses accepted dense output and physical-only error control.

7 Optional JAX execution

An optional JAX driver executes the same numerical map on one selected CPU or GPU, with particle operations batched inside a compiled sequential time loop. It reuses the canonical coefficients and prepared potential splines. Physical-state tolerance scaling, signed stage times, passive energy weights and independent shortened output maps retain their definitions above. Applicable Newton systems use independent particle blocks; an explicitly selected Broyden iteration retains its global secant matrix. Nonlinear iteration limits and failure checks remain active on the device.

Both execution routes use double precision and return the same physical solution layout. Floating-point reassociation can produce round-off differences without changing the designed order or geometric properties. First-call compilation is separate from warmed integration timings, which include transfers and output construction. The sibling runtime document and shared execution contract specify supported callbacks and exact device scope.